



# Prediction the Sale Price of Houses

Springboard Capstone 2 | Swecha Kranthi



# Project Goal

The goal of the project is to use data science tools and statistical analysis to better understand the relationship between house features and how these can be used to predict house prices.

# Stakeholders

- Home Buyers - consumers can make better informed decisions about purchasing a home
- Real Estate Agents - Agents can be more informed about shifting trends in the market
- Companies - companies like Zillow use similar analysis to calculate an estimated value of a house based on its features and surrounding facts.
- Banks - Banks can use similar analysis to assess the risk that comes with a give mortgage on a house

# Why it matters?

Whether you are buying a home or not, Home prices are a key indicator on how well the broader economy is performing. In addition, such analysis can help the key stakeholders to make a more informed decision. This information can help people plan their finances, assess the risk that comes with their mortgage and secure a better financial position for the future.

# Data Information

- The Ames housing data set was compiled by Dr. Dean DeCock of Truman State University. It can be found on Kaggle competition site:  
<https://www.kaggle.com/c/house-prices-advanced-regression-techniques/overview>
- The Data set is comprised of 79 explanatory variables describing (almost) every aspect of residential homes in Ames, Iowa
- There are 2 data sets: Train and Test
  - Train set contains 1460 rows of data with 80 columns. The final column being the Sale Price
  - Test set contain 1459 rows of data with 79 feature columns which we use to predict the final Sale Price.

# Features

- 48 categorical columns
- 32 numerical columns as seen on the right:
- 1 boolean column

Id  
MSSubClass  
LotFrontage  
LotArea  
OverallQual  
OverallCond  
YearBuilt  
YearRemodAdd  
MasVnrArea  
BsmtFinSF1  
BsmtFinSF2  
BsmtUnfSF  
TotalBsmtSF


2ndFlrSF  
LowQualFinSF  
GrLivArea  
BedroomAbvGr  
TotRmsAbvGrd  
GarageYrBlt  
GarageCars  
GarageArea  
WoodDeckSF  
OpenPorchSF  
EnclosedPorch  
3SsnPorch  
ScreenPorch  
PoolArea

PoolArea  
MiscVal  
MoSold  
YrSold  
SalePrice


# Exploratory Data Analysis

Pearson Correlations of features of correlations > 0.5 with Sale Price

```
[11]: Index(['OverallQual', 'YearBuilt', 'YearRemodAdd', 'TotalBsmtSF', '1stFlrSF',
          'GrLivArea', 'FullBath', 'TotRmsAbvGrd', 'GarageCars', 'GarageArea',
          'SalePrice'],
         dtype='object')
```



Pearson Correlations of all features



# Data Cleaning 1

- Identify and Delete that Contain Duplicate Data
  - No duplicates in data
- Address Missing Values
  - We had 19 columns with some missing values. 4 of these columns had more than 50% missing values: Pool, Miscellaneous Features, Alley, Fence. I decided to replace all null values with their respective fill values. All numerical columns were either filled with mean, median or 0 as their fill value. All categorical columns had 'None' as their fill value.

```
#4.1 Delete any duplicate data
duplicate = train_data.duplicated()
print(duplicate.sum())
```

0

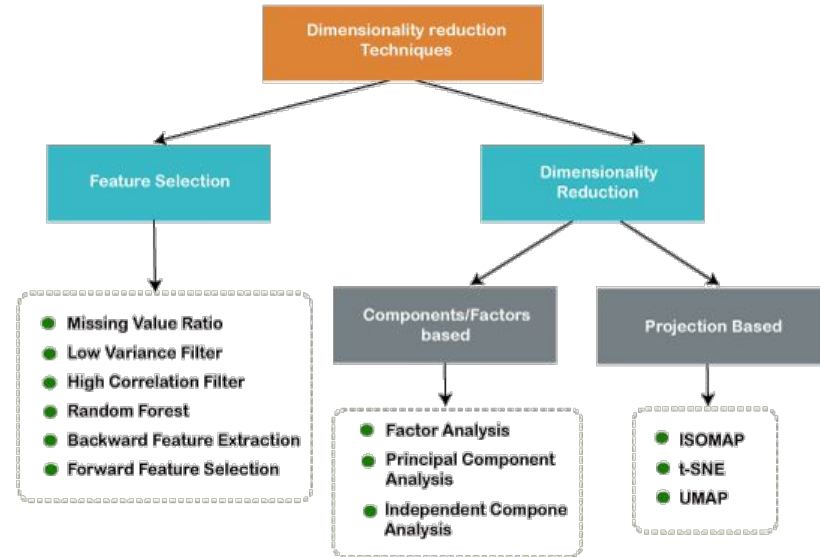
List of null values for each column:

	count	%	types
PoolQC	1453	99.520548	object
MiscFeature	1406	96.301370	object
Alley	1369	93.767123	object
Fence	1179	80.753425	object
FireplaceQu	690	47.260274	object
LotFrontage	259	17.739726	float64
GarageYrBlt	81	5.547945	float64
GarageCond	81	5.547945	object
GarageFinish	81	5.547945	object
GarageQual	81	5.547945	object
GarageType	81	5.547945	object
BsmtFinType2	38	2.602740	object
BsmtExposure	38	2.602740	object
BsmtQual	37	2.534247	object
BsmtCond	37	2.534247	object
BsmtFinType1	37	2.534247	object
MasVnrType	8	0.547945	object
MasVnrArea	8	0.547945	float64
Electrical	1	0.068493	object
Condition1	0	0.000000	object
Neighborhood	0	0.000000	object



# Data Cleaning 2

- Dimensionality Reduction - To reduce the number of variable used to train data
  - There are many ways of doing this, but I used my two favorites: High Correlation Filtering, and Random Forest.
    - High Correlation Filtering is a method where we look at the correlations between each feature and our target value. We decide to keep the once that correlate highly with sale price.
    - Random Forest is a regression algorithm that has built in feature importance, thus we can avoid manually selecting for features and use everything to build our model.



# Preprocessing

- **Numerical Data** needs to be adjusted for any skewness in its distribution. This is to prevent misleading the models and generating wrong predictions.
  - We adjust the skewness of the sale price(right skewed) column by taking the natural log of the data
  - We also adjust the skewness of all numerical columns that have skewness > 0.5 using natural log

```
#Here we take the natural log of skewed data columns to normalize their distributions  
total_data[skew2.index] = np.log1p(total_data[skew2.index])  
total_data
```

- **Categorical Data** is better designed for classification problems. It is hard to use categorical data to creation regressions.
  - To solve this problem, we encode each of our categorical features into a series of numerical features to represent its respective category.

```
total_cat = pd.get_dummies(total_cat)  
total_cat
```

# Modeling 1 - all models

- For Modeling, I used the automation tool - **Pycaret** to test every regression model and find the best one.
- We optimized for lowest Mean Squared Error and High  $R^2$  value
- The result ruled that **Bayesian Ridge Regression** yielded the best model with 88.47% accuracy and the lowest root mean squared error of 0.1216

	Model	MAE	MSE	RMSE	R2	RMSLE	MAPE	TT (Sec)
<b>br</b>	Bayesian Ridge	0.0827	0.0158	0.1216	0.8847	0.0093	0.0069	0.2110
<b>huber</b>	Huber Regressor	0.0876	0.0170	0.1279	0.8777	0.0098	0.0073	0.4680
<b>ridge</b>	Ridge Regression	0.0865	0.0169	0.1263	0.8773	0.0097	0.0072	0.0380
<b>omp</b>	Orthogonal Matching Pursuit	0.0888	0.0176	0.1278	0.8701	0.0098	0.0074	0.0450
<b>gbr</b>	Gradient Boosting Regressor	0.0930	0.0187	0.1341	0.8666	0.0103	0.0078	0.4610
<b>lightgbm</b>	Light Gradient Boosting Machine	0.0945	0.0197	0.1377	0.8595	0.0106	0.0079	0.7940
<b>xgboost</b>	Extreme Gradient Boosting	0.0978	0.0208	0.1423	0.8543	0.0110	0.0082	2.0350
<b>rf</b>	Random Forest Regressor	0.1032	0.0231	0.1503	0.8383	0.0116	0.0086	1.5290
<b>et</b>	Extra Trees Regressor	0.1093	0.0246	0.1553	0.8282	0.0120	0.0091	1.6620
<b>par</b>	Passive Aggressive Regressor	0.1092	0.0238	0.1507	0.8267	0.0116	0.0091	0.0570
<b>knn</b>	K Neighbors Regressor	0.1213	0.0299	0.1717	0.7905	0.0132	0.0101	0.0850
<b>ada</b>	AdaBoost Regressor	0.1288	0.0305	0.1741	0.7864	0.0134	0.0107	0.4050
<b>dt</b>	Decision Tree Regressor	0.1537	0.0475	0.2159	0.6696	0.0167	0.0128	0.0610
<b>lasso</b>	Lasso Regression	0.2976	0.1496	0.3860	-0.0200	0.0295	0.0247	0.0320
<b>en</b>	Elastic Net	0.2976	0.1496	0.3860	-0.0200	0.0295	0.0247	0.0260
<b>llar</b>	Lasso Least Angle Regression	0.2976	0.1496	0.3860	-0.0200	0.0295	0.0247	0.5690
<b>lr</b>	Linear Regression	3.2366	1835.4075	16.0494	-11762.2038	0.1930	0.2691	0.0540

# Modeling 2 - Bayesian Ridge Optimization

- To optimize the BR model, we tested out and tuned the following basic parameters of the BR model:
  - `alpha_init` : Initial value for alpha (precision of the noise). If not set, `alpha_init` is  $1/\text{Var}(y)$ .
  - `lambda_init` : Initial value for lambda (precision of the weights). If not set, `lambda_init` is 1.

```
# Bayesian Ridge Parameters
parameters = {'alpha_init': [1, 1.1, 1.2, 1.3, 1.4, 1.5, 1.9],
              'lambda_init': [1e-1, 1e-2, 1e-3, 1e-4, 1e-5, 1e-6, 1e-9]}
}
```

	MAE	MSE	RMSE	R2	RMSLE	MAPE
0	0.0836	0.0127	0.1127	0.9027	0.0086	0.0070
1	0.0725	0.0096	0.0980	0.9431	0.0076	0.0061
2	0.0731	0.0126	0.1121	0.9120	0.0088	0.0061
3	0.0846	0.0161	0.1270	0.8937	0.0098	0.0071
4	0.1071	0.0246	0.1569	0.8286	0.0122	0.0089
5	0.0995	0.0402	0.2006	0.6393	0.0148	0.0083
6	0.0865	0.0138	0.1173	0.9111	0.0091	0.0073
7	0.0767	0.0101	0.1003	0.9354	0.0076	0.0063
8	0.0784	0.0112	0.1057	0.9373	0.0082	0.0066
9	0.0652	0.0072	0.0850	0.9441	0.0065	0.0054
Mean	0.0827	0.0158	0.1216	0.8847	0.0093	0.0069
SD	0.0121	0.0093	0.0321	0.0880	0.0023	0.0010

	Model	MAE	MSE	RMSE	R2	RMSLE	MAPE
0	Bayesian Ridge	0.0905	0.0292	0.1707	0.8409	0.0131	0.0077

# Predictions

- After finding the optimal Bayesian Ridge Model. WE finalized it and used it to predict the sale price values on the test data. The predictions can be seen on the right.

```
br_final = finalize_model(br_model)
predictions = predict_model(br_final, data = test_data)
```

```
predictions
```

	Id	Label
Aug 16, 2021 2:58 PM		
1	1461	117556.14022900714
2	1462	146048.77543424553
3	1463	177846.63363915216
4	1464	197838.98681724313
5	1465	193371.42728328216
6	1466	171276.75927646863
7	1467	181424.44714079303
8	1468	163740.39680660894
9	1469	186559.4152422597
10	1470	126332.54156807101
11	1471	170959.1774697489
12	1472	100521.953501796
13	1473	98183.25961281336
14	1474	144199.16798159073

# Limitations

- Time
  - The data is sampled between 2006 - 2010; thus, it is not wise to use this model to predict recent home evaluations.
- Location
  - The data represents only 1 city - Ames, Iowa. We cannot use this model to represent housing in other parts of the county.
- Data variance
  - There are a few data columns that have less than 5 unique values. This lack of variance wasn't properly addressing in my cleaning process.
- Low sample size
  - There were 1460 and 1459 samples in Train and Test Respectively. This data is also skewed right. This low sample size means that we cannot assume our data follows a gaussian distributions.
- Outliers
  - I just normalize the distribution of the data. Instead, I should remove the outliers first. This may have led to a decreased accuracy of my model.

# Future Improvements

- In the future, I would love to spend more time creating a cleaner data set. This is my second time properly cleaning large datasets, I could improve my process by more closely understanding the data and optimizing for data skews better.
- This modeling could also be improved by better optimizing for all the parameters that Bayesian Ridge Regression offers.
- I would love to expand this project to larger data sets of sales from Zillow and Trulia detailing all the sales from the last 3 years and include data from the covid years.

# Credits

- Thank you to Raghunandhan Patthar for being an awesome Springboard Mentor.
- Thank you Moez Ali for making such an amazing tool like PyCaret.
- Thank you Dr. Dean DeCock from Truman State University for creating the data set that I was able to work on and learn from.